

# AI Agent Production Deployment & Cost Analysis Guide

A Comprehensive Framework for Deploying LLM-Based Agents at Scale

AI Engineering Team

November 2025  
Version 1.0.0

# Contents

<b>1 Executive Summary</b>	<b>3</b>
1.1 The Core Question . . . . .	3
1.2 Key Findings . . . . .	3
1.2.1 Cost Comparison (100 Concurrent Users, 60,000 queries/month) . . . . .	3
1.3 Decision Rules . . . . .	3
1.3.1 Choose Open-Source LLMs When: . . . . .	3
1.3.2 Choose API-Based LLMs When: . . . . .	3
<b>2 Introduction</b>	<b>5</b>
2.1 What is an LLM-Based Agent? . . . . .	5
2.1.1 Common Use Cases . . . . .	5
2.2 The Cost Challenge . . . . .	5
<b>3 Architecture Patterns</b>	<b>6</b>
3.1 Pattern 1: Self-Hosted Open-Source LLM . . . . .	6
3.2 Pattern 2: API-Based LLM . . . . .	6
3.3 Pattern 3: Hybrid Approach . . . . .	7
<b>4 Infrastructure Options</b>	<b>8</b>
4.1 Cloud Provider Comparison - GPU Instances . . . . .	8
4.2 Model Size vs Hardware Requirements . . . . .	8
<b>5 Cost Models</b>	<b>9</b>
5.1 Model A: Open-Source LLM (Self-Hosted) . . . . .	9
5.1.1 Cost Components . . . . .	9
5.1.2 Example Calculation (100 Users, 7B Model) . . . . .	9
5.2 Model B: API-Based (OpenAI, Anthropic) . . . . .	10
5.2.1 Pricing Structure (OpenAI Example) . . . . .	10
5.2.2 Example Calculation (100 Users) . . . . .	10
<b>6 Cost Calculation Framework</b>	<b>12</b>
6.1 Universal Cost Formula . . . . .	12
6.2 Scaling Calculator . . . . .	12
6.3 Break-Even Analysis . . . . .	13
6.3.1 Break-Even Examples . . . . .	13
<b>7 Quality vs Cost Analysis</b>	<b>14</b>
7.1 Model Quality Comparison . . . . .	14
7.2 Quality-Adjusted Cost Analysis . . . . .	14

<b>8 Scaling Strategies</b>	<b>15</b>
8.1 Horizontal Scaling (More Instances) . . . . .	15
8.2 Vertical Scaling (Bigger Instances) . . . . .	15
8.3 Model Scaling (Tiered Models) . . . . .	16
8.4 Hybrid Scaling (Open-Source + API) . . . . .	16
<b>9 Decision Framework</b>	<b>17</b>
9.1 Decision Tree . . . . .	17
9.2 Common Scenarios . . . . .	17
9.2.1 Scenario 1: Early-Stage Startup (MVP) . . . . .	17
9.2.2 Scenario 2: Healthcare Application . . . . .	18
9.2.3 Scenario 3: Enterprise Customer Support . . . . .	18
<b>10 Best Practices</b>	<b>19</b>
10.1 Infrastructure Best Practices . . . . .	19
10.1.1 Use Reserved Instances . . . . .	19
10.1.2 Implement Caching . . . . .	19
10.1.3 Use Auto-Scaling . . . . .	19
10.2 Cost Optimization Strategies . . . . .	19
10.2.1 Query Batching . . . . .	19
10.2.2 Smart Routing . . . . .	20
10.2.3 Context Window Management . . . . .	20
10.3 Monitoring & Alerting . . . . .	20
<b>11 Conclusion</b>	<b>21</b>
11.1 Key Takeaways . . . . .	21
11.2 Recommended Approach for New Projects . . . . .	21
11.2.1 Phase 1: Start with API (Months 1-3) . . . . .	21
11.2.2 Phase 2: Evaluate Migration (Months 4-6) . . . . .	22
11.2.3 Phase 3: Optimize (Months 7+) . . . . .	22
11.3 Final Recommendations by Organization Size . . . . .	22
11.3.1 Startup (< 50 users) . . . . .	22
11.3.2 Growth Stage (50-500 users) . . . . .	22
11.3.3 Enterprise (> 500 users) . . . . .	22
<b>12 Additional Resources</b>	<b>23</b>
12.1 Cost Calculators . . . . .	23
12.2 Open-Source Models . . . . .	23
12.3 Benchmarks . . . . .	23

# Chapter 1

## Executive Summary

### 1.1 The Core Question

"How much does it cost to run an LLM-based AI agent for production users?"

This guide provides a comprehensive framework for answering this question, applicable to any LLM agent project across different scales, cloud providers, and use cases.

### 1.2 Key Findings

#### 1.2.1 Cost Comparison (100 Concurrent Users, 60,000 queries/month)

Approach	Monthly	Annual	3-Year TCO	Quality	Privacy
Open-Source (GPU)	\$1,500-\$2,500	\$18K-\$30K	\$54K-\$90K	80-90/100	✓
OpenAI GPT-4 API	\$2,000-\$2,500	\$24K-\$30K	\$72K-\$90K	95/100	✗
OpenAI GPT-3.5 API	\$50-\$600	\$600-\$7.2K	\$1.8K-\$21.6K	75-80/100	✗
Hybrid (Both)	\$1,800-\$3,000	\$21.6K-\$36K	\$64.8K-\$108K	85-95/100	~

Table 1.1: Cost comparison across deployment approaches

### 1.3 Decision Rules

#### 1.3.1 Choose Open-Source LLMs When:

- Data privacy is critical (healthcare, finance, government)
- Long-term deployment (> 1 year)
- High query volume (> 50,000/month)
- Need for customization/fine-tuning
- Predictable budget requirements

#### 1.3.2 Choose API-Based LLMs When:

- Quick prototyping or MVP
- Low query volume (< 10,000/month)

- Limited ML/DevOps expertise
- Need for absolute best quality
- Short-term projects (< 6 months)

# Chapter 2

## Introduction

### 2.1 What is an LLM-Based Agent?

An **LLM-based agent** is a software system that uses Large Language Models to:

- Understand natural language queries
- Generate human-like responses
- Perform actions based on instructions
- Maintain conversation context
- Integrate with external systems (RAG, tools, APIs)

#### 2.1.1 Common Use Cases

- Customer support chatbots
- Virtual assistants
- Document Q&A systems
- Code generation tools
- Content creation platforms

### 2.2 The Cost Challenge

Unlike traditional software, LLM agents have **variable operational costs** based on:

- Number of users
- Query frequency
- Model size and quality
- Response length
- Infrastructure choice

This guide helps you **estimate, optimize, and plan** these costs.

# Chapter 3

## Architecture Patterns

### 3.1 Pattern 1: Self-Hosted Open-Source LLM

```
User → Load Balancer → Application Servers (with GPU) → Vector DB  
                                ↓  
                                LLM Model (7B-70B parameters)
```

#### Characteristics:

- Fixed infrastructure costs
- Complete control
- One-time model download
- Requires ML expertise

**Best for:** High-volume, privacy-sensitive applications

### 3.2 Pattern 2: API-Based LLM

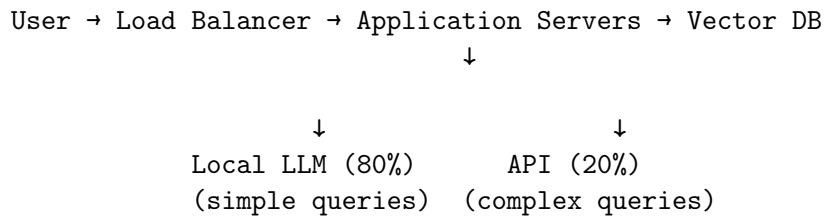
```
User → Load Balancer → Application Servers → Vector DB  
                                ↓  
                                External API (OpenAI, Anthropic)
```

#### Characteristics:

- Pay-per-use pricing
- No infrastructure for LLM
- Instant access to best models
- Variable costs

**Best for:** MVP, low-volume, or quality-critical applications

### 3.3 Pattern 3: Hybrid Approach



#### Characteristics:

- Optimized cost/quality
- Fallback mechanism
- Complexity in routing
- Best of both worlds

**Best for:** Production apps with mixed complexity queries

# Chapter 4

## Infrastructure Options

### 4.1 Cloud Provider Comparison - GPU Instances

Table 4.1: GPU instances for open-source LLMs (Nov 2025 pricing)

Provider	Instance Type	GPU	vCPU	Price/hr	Monthly
AWS	g4dn.xlarge	T4 16GB	4	\$0.526	\$384
AWS	g4dn.2xlarge	T4 16GB	8	\$0.752	\$549
AWS	g5.xlarge	A10G 24GB	4	\$1.006	\$734
Azure	NC4as_T4_v3	T4 16GB	4	\$0.526	\$384
Azure	NC6s_v3	V100 16GB	6	\$3.06	\$2,234
GCP	n1-standard-4+T4	T4 16GB	4	\$0.65	\$475
Alibaba	ecs.gn6v-c8g1.2xlarge	T4 16GB	8	\$0.85	\$621

*Note: Reserved instances offer 30-50% discounts.*

### 4.2 Model Size vs Hardware Requirements

Model Size	Parameters	VRAM	GPU	Speed (tok/s)
Tiny	1-3B	4-8GB	T4, RTX 3060	50-100
Small	7-8B	14-16GB	T4, A10G	30-50
Medium	13-20B	26-40GB	V100, A100	15-30
Large	30-40B	60-80GB	2× A100 40GB	10-20
X-Large	70B+	140GB+	4× A100 40GB	5-10

Table 4.2: Hardware requirements by model size

**Note:** Quantization (8-bit or 4-bit) reduces VRAM by 50-75% with minimal quality loss.

# Chapter 5

## Cost Models

### 5.1 Model A: Open-Source LLM (Self-Hosted)

#### 5.1.1 Cost Components

Total Cost = Infrastructure + Operations + One-Time Setup

##### Infrastructure:

- GPU Instances (scale with users)
- Load Balancer
- Storage (model cache, vector DB)
- Bandwidth
- Backup & Monitoring

##### Operations:

- DevOps Engineer (10-40% time)
- On-call Support
- Model Updates
- Security Audits

#### 5.1.2 Example Calculation (100 Users, 7B Model)

##### Infrastructure (Monthly):

3 × GPU instances (T4 16GB):	$3 \times \$550 = \$1,650$
1 × Vector DB instance:	\$110
Load Balancer:	\$25
Storage (500GB):	\$15
Bandwidth (100 Mbps):	\$60
Backups & Monitoring:	\$40

---

Subtotal: \$1,900

**Operations (Monthly):**

DevOps (20% full-time):	\$500
On-call support:	\$200
Maintenance:	\$100

---

**Subtotal: \$800**

**Total: \$2,700/month = \$32,400/year**

**With 1-year reserved instances (-30%):**

**Total: \$1,890/month = \$22,680/year**

## 5.2 Model B: API-Based (OpenAI, Anthropic)

### 5.2.1 Pricing Structure (OpenAI Example)

Model	Input	Output	Context
GPT-4	\$0.03/1K tok	\$0.06/1K tok	128K
GPT-4 Turbo	\$0.01/1K tok	\$0.03/1K tok	128K
GPT-3.5-Turbo	\$0.0005/1K tok	\$0.0015/1K tok	16K

Table 5.1: OpenAI API pricing (Nov 2025)

### 5.2.2 Example Calculation (100 Users)

#### Usage Assumptions:

- Users: 100 concurrent
- Queries per user per day: 20
- Total monthly queries: 60,000
- Average query: 500 tokens input + 300 tokens output = 800 tokens total
- Monthly tokens: 30M input + 18M output

#### GPT-4 Cost:

Input: $30M \times \$0.03/1K = \$900$
Output: $18M \times \$0.06/1K = \$1,080$

---

API Total: \$1,980/month
Infrastructure: \$300/month
Operations: \$200/month

---

**Total: \$2,480/month = \$29,760/year**

**GPT-3.5-Turbo Cost:**

Input:  $30M \times \$0.0005/1K = \$15$

Output:  $18M \times \$0.0015/1K = \$27$

---

API Total: \$42/month

Infrastructure: \$300/month

Operations: \$200/month

---

**Total:** \$542/month = \$6,504/year

# Chapter 6

## Cost Calculation Framework

### 6.1 Universal Cost Formula

```
1 def calculate_monthly_cost(
2     users: int,
3     queries_per_user_per_day: int,
4     avg_tokens_per_query: int,
5     deployment_type: str, # "open_source" or "api"
6     model_type: str,      # "7B", "13B", "70B" or "gpt-4"
7     cloud_provider: str   # "aws", "azure", "gcp", "alibaba"
8 ) -> float:
9
10    total_monthly_queries = users * queries_per_user_per_day * 30
11
12    if deployment_type == "open_source":
13        gpu_instances = calculate_gpu_instances(users, model_type)
14        hourly_cost = get_gpu_hourly_cost(cloud_provider, model_type)
15
16        infrastructure = gpu_instances * hourly_cost * 730
17        operations = 500 + (gpu_instances * 200)
18
19        return infrastructure + operations
20
21    elif deployment_type == "api":
22        monthly_tokens = total_monthly_queries * avg_tokens_per_query
23        api_cost = calculate_api_cost(model_type, monthly_tokens)
24        infrastructure = 300
25        operations = 200
26
27        return api_cost + infrastructure + operations
```

Listing 6.1: Cost calculation function

### 6.2 Scaling Calculator

Table 6.1: Cost scaling by user count

Users	Queries/Mo	GPUs	Open-Source	GPT-4	GPT-3.5
10	6,000	1	\$850	\$248	\$54
50	30,000	2	\$1,700	\$1,240	\$271
100	60,000	3	\$2,700	\$2,480	\$542

Users	Queries/Mo	GPUs	Open-Source	GPT-4	GPT-3.5
500	300,000	8	\$8,500	\$12,400	\$2,710
1,000	600,000	15	\$15,500	\$24,800	\$5,420
10,000	6,000,000	100	\$110,000	\$248,000	\$54,200

**Key Insight:** Open-source becomes dramatically cheaper at scale ( $> 500$  users).

## 6.3 Break-Even Analysis

The break-even point is when open-source infrastructure costs equal API costs:

$$\text{Break-even} = \frac{\text{Infrastructure Fixed Cost}}{\text{Monthly Savings}}$$

### 6.3.1 Break-Even Examples

- **At 100 users:** Open-source (\$2,700) vs GPT-4 (\$2,480) - GPT-4 slightly cheaper
- **At 500 users:** Open-source (\$8,500) vs GPT-4 (\$12,400) - **46% savings**
- **At 1,000 users:** Open-source (\$15,500) vs GPT-4 (\$24,800) - **60% savings**

#### Rule of Thumb:

- $< 50$  users: Use APIs (lower entry cost)
- 50-100 users: Break-even zone (depends on usage)
- $> 100$  users: Open-source becomes economical
- $> 500$  users: Open-source is significantly cheaper

# Chapter 7

## Quality vs Cost Analysis

### 7.1 Model Quality Comparison

Table 7.1: Benchmark scores across model types

Model	Overall	Reason	Code	Math	Creative	Cost	Tier
GPT-4	95/100	Exc	Exc	Exc	Exc	High	
Claude 3 Opus	94/100	Exc	Exc	Exc	Exc	High	
GPT-4 Turbo	93/100	Exc	Exc	Exc	V.Good	Medium	
Claude 3 Sonnet	88/100	V.Good	V.Good	V.Good	V.Good	Medium	
Mixtral 8x7B	87/100	V.Good	V.Good	Good	V.Good	Low	
Llama 3 70B	86/100	V.Good	V.Good	Good	V.Good	Low	
Mistral 7B	84/100	Good	Good	Good	Good	V.Low	
Llama 3 8B	82/100	Good	Good	Fair	Good	V.Low	
GPT-3.5-Turbo	78/100	Good	Good	Fair	Good	V.Low	
Phi-3-mini	72/100	Fair	Fair	Fair	Fair	Ultra Low	

### 7.2 Quality-Adjusted Cost Analysis

Cost per Quality Point (Annual, 100 users):

GPT-4:  $\$30,000/95 = \$316$  per quality point

Claude Opus:  $\$32,000/94 = \$340$  per quality point

Mixtral 8x7B:  $\$28,000/87 = \$322$  per quality point

Mistral 7B:  $\$23,000/84 = \$274$  per quality point ✓ **Best value**

GPT-3.5-Turbo:  $\$6,500/78 = \$83$  per quality point ✓ **Ultra budget**

# Chapter 8

## Scaling Strategies

### 8.1 Horizontal Scaling (More Instances)

**Approach:** Add more GPU servers as users increase

User Range	GPU Instances
1-50 users	1 instance
50-150 users	2 instances
150-300 users	3 instances
300-500 users	5 instances
500-1000 users	8-10 instances

Table 8.1: Horizontal scaling guidelines

#### Pros:

- Predictable scaling
- High availability (redundancy)
- Easy load balancing

#### Cons:

- Higher fixed costs
- More complex deployment
- Increased operational overhead

### 8.2 Vertical Scaling (Bigger Instances)

**Approach:** Upgrade to more powerful GPUs

#### Pros:

- Simpler architecture
- Lower network overhead
- Can handle larger models

#### Cons:

GPU Type	Monthly Cost
T4 16GB	\$550
V100 32GB	\$2,200
A100 40GB	\$3,000
A100 80GB	\$5,000

Table 8.2: Vertical scaling options

- Single point of failure
- Limited by hardware ceiling
- More expensive per user at low scale

### 8.3 Model Scaling (Tiered Models)

**Approach:** Route queries to appropriate model based on complexity

```

1 def route_query(query, complexity):
2     if complexity == "simple":
3         return small_model_7b    # 70% of queries
4     elif complexity == "medium":
5         return medium_model_13b # 20% of queries
6     else:
7         return large_model_70b # 10% of queries

```

Listing 8.1: Multi-model routing

#### Benefits:

- Optimize cost by using smaller models when possible
- Reserve expensive models for complex queries
- 30-50% cost reduction with minimal quality impact

### 8.4 Hybrid Scaling (Open-Source + API)

**Approach:** Use local models for most queries, API for exceptions

```

1 def handle_query(query):
2     # Try local model first
3     response = local_model.generate(query)
4
5     # Check confidence score
6     if confidence(response) < 0.7:
7         # Fallback to API for better quality
8         response = api_model.generate(query)
9
10    return response

```

Listing 8.2: Hybrid fallback strategy

#### Cost Impact:

- 80% queries → Local model: \$2,000/month
- 20% queries → API fallback: \$500/month
- Total: \$2,500/month
- Benefit: Best quality + acceptable cost

# Chapter 9

## Decision Framework

### 9.1 Decision Tree

1. **Is data privacy critical?** (healthcare, finance, government)

- YES → Use Open-Source (mandatory for compliance)
- NO → Continue

2. **Query volume per month?**

- < 10,000 → Use API (lower entry cost)
- 10,000-50,000 → Evaluate both options
- > 50,000 → Use Open-Source (better economics)

3. **Project timeline?**

- < 3 months → Use API (faster to market)
- 3-12 months → Hybrid (start API, migrate)
- > 12 months → Use Open-Source (long-term savings)

4. **Quality requirements?**

- Must be best → Use GPT-4/Claude
- Very good is okay → Use Mistral/Llama
- Good enough → Use GPT-3.5

5. **Team expertise?**

- No ML experience → Use API (easier)
- Some ML knowledge → Hybrid approach
- Strong ML team → Use Open-Source

### 9.2 Common Scenarios

#### 9.2.1 Scenario 1: Early-Stage Startup (MVP)

- **Users:** 50
- **Budget:** \$5,000/month

- **Timeline:** 3 months
- **Team:** 2 engineers (no ML experience)

**Recommendation:** GPT-3.5-Turbo API

**Cost:** ~\$300/month

**Rationale:** Fast to market, minimal setup, low risk

### 9.2.2 Scenario 2: Healthcare Application

- **Users:** 200
- **Budget:** \$15,000/month
- **Data:** HIPAA-compliant
- **Timeline:** 6 months
- **Team:** 5 engineers (1 ML expert)

**Recommendation:** Self-hosted Llama 3 70B

**Cost:** ~\$5,000/month

**Rationale:** Data privacy mandatory, sufficient budget, long-term project

### 9.2.3 Scenario 3: Enterprise Customer Support

- **Users:** 1,000
- **Budget:** \$50,000/month
- **Quality:** High priority
- **Timeline:** 12 months
- **Team:** 10 engineers (3 ML experts)

**Recommendation:** Hybrid (Mistral 7B + GPT-4 fallback)

**Cost:** ~\$18,000/month

**Rationale:** Cost-effective at scale, best quality when needed

# Chapter 10

## Best Practices

### 10.1 Infrastructure Best Practices

#### 10.1.1 Use Reserved Instances

- **Savings:** 30-50% for 1-3 year commitments
- **Example:** \$2,000/month → \$1,400 (1-year) → \$1,000 (3-year)
- **When:** Predictable usage for > 6 months

#### 10.1.2 Implement Caching

```
1 # Cache frequent queries
2 cache_hit_rate = 0.30-0.40  # Typical for customer support
3 cost_reduction = cache_hit_rate * compute_cost
4 # Example: 40% hit rate = 40% reduction in compute
```

Listing 10.1: Response caching

#### 10.1.3 Use Auto-Scaling

```
1 min_instances: 2
2 max_instances: 10
3 scale_up_threshold: 75% GPU utilization
4 scale_down_threshold: 30% GPU utilization
```

Listing 10.2: Auto-scaling configuration

### 10.2 Cost Optimization Strategies

#### 10.2.1 Query Batching

- Single query: 100ms compute
- Batched (10 queries): 150ms compute
- Efficiency: 15ms per query (85% savings)

### 10.2.2 Smart Routing

```

1 def route_query(query):
2     if is_faq(query):
3         return cached_response # $0
4     elif is_simple(query):
5         return small_model      # $0.001
6     else:
7         return large_model    # $0.01

```

Listing 10.3: Cost-aware routing

### 10.2.3 Context Window Management

```

1 # Don't send entire conversation history
2 max_context_tokens = 2000 # vs 8000
3 cost_reduction = 0.75 # 75% reduction

```

Listing 10.4: Optimize context usage

## 10.3 Monitoring & Alerting

### Key Metrics to Track:

1. Cost per query
2. Average response time
3. Model quality (user satisfaction)
4. GPU utilization
5. Cache hit rate
6. Error rate
7. API quota usage (if hybrid)

### Alert Thresholds:

- Cost per query >\$0.05 → investigate inefficiency
- GPU utilization > 85% → scale up
- GPU utilization < 30% for 2h → scale down
- Error rate > 1% → check model health
- API rate limit > 80% → upgrade tier

# Chapter 11

## Conclusion

### 11.1 Key Takeaways

#### 1. No One-Size-Fits-All Solution

- Small scale (< 100 users): APIs often cheaper
- Large scale (> 500 users): Self-hosted more economical
- Hybrid: Best quality-cost balance for many scenarios

#### 2. Cost Scales Differently

- APIs: Linear with usage (predictable per query)
- Self-hosted: Stepped with scale (fixed + marginal)
- Crossover point: 50-200 users typically

#### 3. Quality vs Cost Trade-off

- Premium models (GPT-4): 20% better, 2-4× more expensive
- Mid-tier (Mistral): 90% quality, 30-50% cheaper
- Budget (GPT-3.5): 80% quality, 90% cheaper

#### 4. Non-Cost Factors Often Decide

- Data privacy requirements
- Compliance needs
- Team expertise
- Time to market

### 11.2 Recommended Approach for New Projects

#### 11.2.1 Phase 1: Start with API (Months 1-3)

- Quick validation
- Learn usage patterns
- Collect data for potential fine-tuning
- **Cost:** \$500-2,000/month

### 11.2.2 Phase 2: Evaluate Migration (Months 4-6)

- Calculate break-even point
- Test open-source models
- Compare quality metrics
- **Decision:** Migrate or stay with API

### 11.2.3 Phase 3: Optimize (Months 7+)

- Fine-tune models on your data
- Implement hybrid if beneficial
- Continuous cost monitoring
- **Cost:** Optimized for your scale

## 11.3 Final Recommendations by Organization Size

### 11.3.1 Startup (< 50 users)

- **Approach:** GPT-3.5-Turbo or Claude Haiku
- **Cost:** \$300-500/month
- **Focus:** Product-market fit, not infrastructure

### 11.3.2 Growth Stage (50-500 users)

- **Approach:** Hybrid (self-hosted + API fallback)
- **Cost:** \$2,000-8,000/month
- **Focus:** Balance cost optimization with quality

### 11.3.3 Enterprise (> 500 users)

- **Approach:** Self-hosted with fine-tuning
- **Cost:** \$8,000-50,000/month
- **Focus:** Full control, best economics at scale

# Chapter 12

## Additional Resources

### 12.1 Cost Calculators

- AWS Pricing Calculator: <https://calculator.aws>
- Azure Pricing Calculator: <https://azure.microsoft.com/pricing/calculator>
- GCP Pricing Calculator: <https://cloud.google.com/products/calculator>
- OpenAI API Pricing: <https://openai.com/pricing>
- Anthropic Pricing: <https://anthropic.com/pricing>

### 12.2 Open-Source Models

- Hugging Face Model Hub: <https://huggingface.co/models>
- Ollama (local deployment): <https://ollama.ai>
- vLLM (efficient serving): <https://github.com/vllm-project/vllm>

### 12.3 Benchmarks

- LLM Leaderboard: [https://huggingface.co/spaces/HuggingFaceH4/open\\_llm\\_leaderboard](https://huggingface.co/spaces/HuggingFaceH4/open_llm_leaderboard)
- Chatbot Arena: <https://chat.lmsys.org>

---

**Remember:** The best solution depends on your specific requirements.  
Use this framework as a starting point, but always validate with your own testing.

Good luck with your LLM agent deployment!

*Last Updated: November 2025  
Version 1.0.0*